

ECE60146 Homework 2 Report

Lucas Nguyen <nguye800@purdue.edu>

****There seems to be a mismatch in instructions as the instructions in 3.3 seem to be similar to the Task Summary in 3.6 but with different numbers so I will just carry out the instructions laid out in 3.6 since the skeleton code is more oriented towards those instructions****

Introduction

The objective of this homework is to get familiar with datasets, dataloaders, and transformations. The assignments within this homework include creating datasets and using dataloaders to output the dataset in batches with various transformations.

Creating the Environment

The preliminary step for this homework is to create a conda environment which contains all packages required to complete this homework (except for matplotlib.pyplot which needs to be downloaded). The steps taken to create the environment is as follows:

1. Sudo pip install conda
 - a. Installs conda so that a virtual environment can be created
2. conda create -name ece60146 python=3.11
 - a. creates a new environment called ece60146 with python 3.11
3. conda activate ece60146
 - a. Activates the newly created environment. Now any commands run in the terminal will be using this environment
4. Conda install pytorch torchvision -c pytorch
 - a. Within the new environment pytorch and torchvision are installed which are dependencies for this homework
5. Conda env export > environment.yml
 - a. Exports environment as a yaml file for grading

1 Load ImageNet (Option A)

Implementing a dataset to load the ImageNet dataset using the zip file from Brightspace (Option A). The `_load_samples` function loops through each class label given in initialization and appends `images_per_class` images for each class label to `self.samples`. The assignment was to load 50 images total (5 from each class) so a `ImageNetSubset` instantiation was initialized with `images_per_class=5` and batches were looped through 10 times in a for loop to give a total of 50 images. Code output proves that 5 images from each label class were correctly output.

`_load_samples` Implementation

```
def _load_samples(self):
    """Load image paths for each requested class."""
    for label in self.class_labels:
        base_path = os.path.join(self.root, str(label))
        count = 0
        for img in sorted(os.listdir(base_path)):
            if img.lower().endswith('.png', '.jpg', '.jpeg'):
                self.samples.append((os.path.join(base_path, img), label))
                count += 1

            if count == self.images_per_class:
                break
```

`__getitem__` Implementation

```
def __getitem__(self, index):
    """
    Returns:
        image (Tensor): Transformed image
        label (int): Class label
    """

    img_path = self.samples[index][0]
    image = Image.open(img_path).convert("RGB")
    if self.transform:
        image = self.transform(image)

    return image, self.samples[index][1]
```

Load 50 images (5 per class) Code Implementation

```
217     # Task 1: Load ImageNet (50 images: 10 classes x 5 images)
218     print("\n" + "=" * 60)
219     print("Task 1: Loading ImageNet")
220     print("=" * 60)
221
222     imagenet = ImageNetSubset(root='D:\Lucas College\Purdue\Y4\ECE60146-HW\HW2\imagenet_subset',
223                               class_labels=required_labels,
224                               images_per_class=5,
225                               transform=transform_basic)
226
227     dataloader = DataLoader(
228         imagenet,
229         batch_size=5,
230         shuffle=False,
231         num_workers=4,
232         pin_memory=True
233     )
234
235     it = iter(dataloader)
236     print("Total Images:", len(imagenet))
237     for i in range(10):
238         images, labels = next(it)
239         print("Class Label:", labels)
```

Code Output

```
=====
Task 1: Loading ImageNet
=====
Loaded 50 images from 10 classes
Total Images: 50
Class Label: tensor([1, 1, 1, 1, 1])
Class Label: tensor([151, 151, 151, 151, 151])
Class Label: tensor([281, 281, 281, 281, 281])
Class Label: tensor([291, 291, 291, 291, 291])
Class Label: tensor([325, 325, 325, 325, 325])
Class Label: tensor([386, 386, 386, 386, 386])
Class Label: tensor([430, 430, 430, 430, 430])
Class Label: tensor([466, 466, 466, 466, 466])
Class Label: tensor([496, 496, 496, 496, 496])
Class Label: tensor([950, 950, 950, 950, 950])
```


2 Visualize ImageNet

This assignment entailed loading 10 images into a 2x5 grid (1 per class) therefore ImageNetSubset was initialized with images_per_class=1 and the DataLoader was set to batch_size=10 to get all 10 images at once then plotted in a for loop from 0-9 inclusive.

Code Implementation

```
245 # Task 2: Visualize ImageNet (1 image per class = 10 images
246 print("\n" + "=" * 60)
247 print("Task 2: Visualizing ImageNet")
248 print("=" * 60)
249
250 imagenet = ImageNetSubset(root='D:\Lucas College\Purdue\Y4\ECE60146-HW\HW2\imagenet_subset',
251                           class_labels=required_labels,
252                           images_per_class=1,
253                           transform=transform_basic)
254
255 dataloader = DataLoader([
256     imagenet,
257     batch_size=10,
258     shuffle=False,
259     num_workers=4,
260     pin_memory=True
261 ])
262
263 images, labels = next(iter(dataloader))
264 fig, axes = plt.subplots(2, 5, figsize=(12, 4))
265
266 for i in range(10):
267     img = images[i].permute(1, 2, 0) # CHW → HWC
268     axes[i//5][i%5].imshow(img)
269     axes[i//5][i%5].axis('off')
270     axes[i//5][i%5].set_title(f"Image {labels[i]}")
271
272 plt.tight_layout()
273 plt.savefig("imagenet_samples.png", dpi=300, bbox_inches="tight")
274 plt.show()
```

Output (Image #: Class number listed in folder)




3 Augmentation Comparison

This section required plotting 3 images alongside their augmented versions. Two ImageNetSubsets were initialized, one with transform_basic to plot the original image, and one with transform_custom to show the augmented version. Then two dataloaders were initialized for the two datasets with shuffle=False so that they would both get the same image in the same order in order to properly plot them against each other.

Augmentation Code

```
276 # Task 3: Show augmentation effects
277 print("\n" + "=" * 60)
278 print("Task 3: Augmentation comparison")
279 print("=" * 60)
280
281 imangenet_orig = ImageNetSubset(root='D:\Lucas College\Purdue\Y4\ECE60146-HW\HW2\imagenet_subset',
282                                 class_labels=required_labels,
283                                 images_per_class=1,
284                                 transform=transform_basic)
285 imangenet_aug = ImageNetSubset(root='D:\Lucas College\Purdue\Y4\ECE60146-HW\HW2\imagenet_subset',
286                                 class_labels=required_labels,
287                                 images_per_class=1,
288                                 transform=transform_custom)
289
290 dataloader_orig = DataLoader(
291     imangenet_orig,
292     batch_size=3,
293     shuffle=False,
294     num_workers=4,
295     pin_memory=True
296 )
297
298 dataloader_aug = DataLoader(
299     imangenet_aug,
300     batch_size=3,
301     shuffle=False,
302     num_workers=4,
303     pin_memory=True
304 )
305
306 images_orig, labels_orig = next(iter(dataloader_orig))
307 images_aug, labels_aug = next(iter(dataloader_aug))
308 fig, axes = plt.subplots(2, 3, figsize=(12, 14))
309
310 for i in range(3):
311     img_orig = images_orig[i].permute(1, 2, 0) # CHW → HWC
312     axes[0][i].imshow(img_orig)
313     axes[0][i].axis('off')
314     axes[0][i].set_title(f'Image {labels_orig[i]}')
315
316     img_aug = images_aug[i].permute(1, 2, 0) # CHW → HWC
317     axes[1][i].imshow(img_aug)
318     axes[1][i].axis('off')
319     axes[1][i].set_title(f'Image Aug [{labels_aug[i]}]')
320
321 plt.tight_layout()
322 plt.savefig("imagenet_aug.png", dpi=300, bbox_inches="tight")
323 plt.show()
```

Output



4 Custom Dataset

This section required implementing a custom dataset class where images from a given root path would be used to form the dataset. `self.image_paths` was implemented to loop through the root directory and build an array of every image type file. `__getitem__` would convert the image to RGB then apply the desired transform before returning the image and a default label 0.


Class Implementation

```
134 # =====
135 # Custom Dataset for your own images
136 # =====
137
138 class CustomDataset(Dataset):
139     """Dataset for loading custom images from a folder."""
140
141     def __init__(self, root, transform=None):
142         self.root = root
143         self.transform = transform
144
145         self.image_paths = [os.path.join(root, img) for img in os.listdir(root) if img.lower().endswith('.png', '.jpg', '.jpeg', '.bmp', '.gif')]
146
147         print(f"Found {len(self.image_paths)} images in {root}")
148
149     def __len__(self):
150         return len(self.image_paths)
151
152     def __getitem__(self, index):
153         img_path = self.image_paths[index]
154         image = Image.open(img_path).convert("RGB")
155         if self.transform:
156             image = self.transform(image)
157
158         return image, 0
```

Code Implementation

```
325 # Task 4: Custom dataset
326 print("\n" + "=" * 60)
327 print("Task 4: Custom dataset")
328 print("=" * 60)
329
330 fig, axes = plt.subplots(2, 5, figsize=(12, 4))
331
332 customDataset_1 = CustomDataset(root='D:/Lucas College/Purdue/Y4/ECE60146-HW/HW2/imagenet_subset/151', transform=transform_custom)
333 customDataset_2 = CustomDataset(root='D:/Lucas College/Purdue/Y4/ECE60146-HW/HW2/imagenet_subset/281', transform=transform_custom)
334
335 dataset20 = ConcatDataset([customDataset_1, customDataset_2])
336
337 dataloader_custom = DataLoader(
338     dataset20,
339     batch_size=10,
340     shuffle=True,
341     num_workers=4,
342     pin_memory=True
343 )
344 num_generated = 0
345 for epoch in range(5):
346     images, labels = next(iter(dataloader_custom))
347     for i in range(10):
348         num_generated += 1
349         if epoch == 4:
350             img = images[i].permute(1, 2, 0) # CHW → HWC
351             axes[1//5][i%5].imshow(img)
352             axes[1//5][i%5].axis("off")
353             axes[1//5][i%5].set_title(f"Image {labels[i]}")
354
355 print("Number of Augmented Images Generated:", num_generated)
356 plt.tight_layout()
357 plt.savefig("custom_samples.png", dpi=300, bbox_inches="tight")
358 plt.show()
```

Code Output



5 DataLoader Performance

This step involved initializing a dataloader with different batch sizes and workers then timing how long it takes to load + transform + batch over one pass of the dataloader. From the numbers shown in the table the most likely reasoning is that the transform isn't computationally heavy enough to add significant time to augment images and thus the difference between num_workers=0 with different batch sizes is most likely either due to noise or to the smaller time required to internally batch the entire dataset in batches of 16 vs. 64. One major difference is that the two tests with num_workers=4 took significantly longer than the other tests. One reason for this is that the overhead to initialize several concurrent workers is much greater than the time the augment the images and therefore adds roughly a ~2.5 sec increase.

Code Implementation

```
356     # Task 5: DataLoader performance
357     print("\n" + "=" * 60)
358     print("Task 5: DataLoader performance")
359     print("=" * 60)
360
361     imagenet = ImageNetSubset(root='D:\\Lucas College\\Purdue\\Y4\\ECE60146-HW\\HW2\\imagenet_subset',
362                               class_labels=required_labels,
363                               images_per_class=10,
364                               transform=transform_custom)
365
366     for pair in [(16 ,0), (16 ,4), (64 ,0), (64 ,4)]:
367         t0 = time.perf_counter()
368         dataloader = DataLoader(
369             imagenet,
370             batch_size=pair[0],
371             shuffle=False,
372             num_workers=pair[1],
373             pin_memory=True
374         )
375         seen = 0
376         for images, labels in dataloader:
377             seen += images.size(0)
378         t1 = time.perf_counter()
379         print(f"For batch size: {pair[0]} and num workers: {pair[1]} the time was {t1-t0}")
```

Code Output

```
=====
Task 5: DataLoader performance
=====
Loaded 100 images from 10 classes
For batch size: 16 and num workers: 0 the time was 0.2791416998952627
For batch size: 16 and num workers: 4 the time was 2.930186300072819
For batch size: 64 and num workers: 0 the time was 0.26909289974719286
For batch size: 64 and num workers: 4 the time was 2.936049100011587
```

Performance Table

Batch_size	Num_workers	Time(sec)
16	0	0.40694050025194883
16	4	2.946450099814683
64	0	0.39657749980688095
64	4	2.9996744003146887

6 RGB Statistics

This section involved displaying the difference in values before and after the normalize transform. Before any transforms the min and max values for the pixels were 0 and 255 which makes sense because 8 bit rgb values can only range between 0-255. Secondly the values were taken after the transform_basic which just resizes and uses ToTensor. The value range makes sense because ToTensor transforms data to range from 0-1. Lastly the transform_custom was performed on the data which applies the normalize function to the data. The normalize function given had slightly different mean and std for each channel but overall normalized the RGB channels to roughly [-2.1, 2.6].

Code Implementation

```
381 # Task 6: RGB statistics
382 print("\n" + "=" * 60)
383 print("Task 6: RGB statistics")
384 print("=" * 60)
385
386 def channel_min_max(x):
387     return [(x[c].min().item(), x[c].max().item()) for c in range(3)]
388
389 img = Image.open('D:/Lucas College/Purdue/Y4/ECE60146-HW/Hw2/imagenet_subset/1/00001.jpeg').convert("RGB")
390 img_pre = transform_basic(img)
391 img_post = transform_custom(img)
392
393 img_stats = channel_min_max(np.array(img))
394 pre_stats = channel_min_max(img_pre)
395 post_stats = channel_min_max(img_post)
396
397 print("Before all transformation")
398 print("R:", img_stats[0], "G:", img_stats[1], "B:", img_stats[2])
399
400 print("After ToTensor before Normalize")
401 print("R:", pre_stats[0], "G:", pre_stats[1], "B:", pre_stats[2])
402
403 print("After Normalize")
404 print("R:", post_stats[0], "G:", post_stats[1], "B:", post_stats[2])
```

Code Output

```
=====
Task 6: RGB statistics
=====
Before all transformation
R: (0, 255) G: (0, 255) B: (0, 255)
After ToTensor before Normalize
R: (0.0, 0.9960784316062927) G: (0.0, 0.9960784316062927) B: (0.003921568859368563, 0.9960784316062927)
After Normalize
R: (-2.1179039478302, 2.231783390045166) G: (-2.0357141494750977, 2.393557548522949) B: (-1.804444432258606, 2.6051416397094727)
```

7 Reproducibility

This step was to demonstrate the effect of using random seeds on batch ordering. If no seed is specified, pytorch initializes its RNG from system entropy which explains the difference in images between the first 2 runs. The second 2 runs were initialized with the same seed meaning that even though shuffle=True, the “randomness” of the batch order is actually the same because the seed is the same, hence why the last 2 runs have the same images.

Code Implementation

```
406     # Task 7: Reproducibility
407     print("\n" + "=" * 60)
408     print("Task 7: Reproducibility")
409     print("=" * 60)
410
411     def set_seed(seed = 60146):
412         random.seed(seed)
413         np.random.seed(seed)
414         torch.manual_seed(seed)
415         if torch.cuda.is_available():
416             torch.cuda.manual_seed_all(seed)
417
418     set_seed()
419
420     imagenet = ImageNetSubset(root='D:\Lucas College\Purdue\Y4\ECE60146-HW\HW2\imagenet_subset',
421                             class_labels=required_labels,
422                             images_per_class=10,
423                             transform=transform_custom)
424
425     dataloader = DataLoader(
426         imagenet,
427         batch_size=2,
428         shuffle=True,
429         num_workers=0,
430         pin_memory=True
431     )
432
433     images, labels = next(iter(dataloader))
434     fig, axes = plt.subplots(1, 2, figsize=(12, 4))
435
436     for i in range(2):
437         img = images[i].permute(1, 2, 0) # CHW -> HWC
438         axes[i].imshow(img)
439         axes[i].axis('off')
440         axes[i].set_title(f"Image {labels[i]}")
441
442     plt.tight_layout()
443     plt.savefig("imagenet_samples.png", dpi=300, bbox_inches="tight")
444     plt.show()
```

No Seed First Run



No Seed Second Run



Seed 60146 First Run



Seed 60146 Second Run

